

How Does AI Audio Generation Work?

Caden Khuu

1 Introduction

In the rapidly developing field of AI technologies, one of the most recent and intriguing applications is the generation of music and speech. Services such as Uberduck have taken the internet by storm with the seemingly miraculous ability to generate any phrase the user desires in the voice of cartoon characters or celebrities. Jukebox AI has made it possible to create "extensions" of songs given snippets less than 20 seconds in length. At the core of AI music, audio, and speech generation is the utilization of neural networks, which are computational models inspired by the human brain's interconnected neurons. These networks are trained on vast datasets of existing musical compositions, spoken words, or diverse audio samples, enabling them to learn complex patterns and structures inherent in the data. As the models learn, they become capable of generating new, original content that often exhibits a remarkable level of creativity and coherence. In the realm of music generation, AI systems can compose melodies, harmonies, and rhythms across various genres. They can analyze the styles of different composers and create music that emulates their unique characteristics or blend influences to produce entirely novel compositions. Similarly, in the field of audio synthesis, AI algorithms can recreate lifelike sounds or generate entirely new ones, providing unprecedented possibilities for sound design in industries such as gaming, film, and virtual reality. The application of AI in speech generation has yielded systems that can produce natural-sounding human speech, often indistinguishable from recordings of real voices. These technologies find applications in voice assistants, audiobooks, and voiceovers, contributing to enhanced user experiences across various platforms.

2 History of Development

The evolution of AI music generation traces its roots back to the mid-20th century when researchers began exploring algorithmic composition as a means of creating music through mathematical rules. Pioneering works such as the 1957 composition ILLIAC Suite, the first string quartet composed entirely on a computer program, laid the foundation for early experiments in computer-generated music. The advent of MIDI, a standard file format to store music, in the 1980s marked a significant milestone, allowing musicians to transmit musical ideas digitally with a number of virtual instruments and paving the way for digital synthesis. In the 1990s, a new wave of innovation emerged with the application of evolutionary algorithms and genetic programming to music composition. This period saw the development of interactive installations like "Emily Howell," an early attempt at AI-generated classical-style music by David Cope. These endeavors hinted at the potential of AI in the creative realm. The 2010s witnessed a paradigm shift with the rise of neural networks and deep learning. Recurrent Neural Networks (RNNs) and Long Short-Term Memory (LSTM) networks became popular for sequence generation tasks, including music. Google's Magenta project, launched in 2016, played a pivotal role by providing tools and models for AI music generation, making the technology more accessible to both researchers and artists. The integration

of AI-generated music into popular culture gained momentum during this period. Platforms like AIVA and Jukedeck offered AI-generated music for commercial use, showcasing the technology's potential in the mainstream music scene. As we entered the 2020s, transformer models, exemplified by OpenAI's GPT, revolutionized natural language processing and, consequently, music generation. These models could understand and generate more complex musical structures, expanding the possibilities of AI-generated compositions. Style transfer techniques and hybrid models combining rule-based approaches with deep learning further enriched the expressiveness of AI-generated music, allowing the mimicry of specific genres or artists. Real-time collaboration between human musicians and AI systems became a focal point of innovation, exemplified by platforms like Amper Music and IBM's Watson Beat. These tools demonstrated the potential for collaborative creativity, where AI serves as a dynamic and responsive partner in the creative process. AI-generated music also found applications in education, assisting learners in composition, arrangement, and exploration of diverse musical styles. The continuous interplay between technological advancements, artistic exploration, and the quest to push the boundaries of musical creativity defines the rich history of AI music generation. As AI continues to evolve, its role in shaping the future of music composition and performance promises to be even more transformative.

3 How it Works

One of the most popular services for AI generated audio is Uberduck, a popular online service that allows users to effortlessly convert text into speech by selecting from a diverse range of voices, including those of cartoon characters and celebrities. The platform is just one of many that utilize AI vocal imitation and voice cloning techniques, a process wherein an extensive database of a specific character's voice serves as the foundation. The training data of audio clips of the voice is inputted into the system and a neural network learns to mimic the patterns, tone, and nuances of the voice by analyzing the acoustic features, intonation, and speech patterns present in the given audio. In this case, Uberduck uses a transformer model architecture, commonly used for natural language processing, to analyze the text prompt and convert it into audio format. Using similar techniques to this, machine learning developers have been able to create a wide array of AI-assisted audio manipulations, from editing the pitch so that an AI voice could sing a song, to being able to record your own speech and playing it back with a different tone and accent. In addition to human voices, AI is also capable of generating music. Jukebox AI, exemplified by models such as OpenAI's implementation, revolutionize music generation through its advanced generation of diverse musical genres and styles. At its core, a user inputs a text prompt of what they wish the outputted song to sound like, and the system begins by meticulously collecting and analyzing extensive datasets encompassing a wide array of musical compositions. The training data is encoded by a generative mode called VQ-VAE, which compresses the audio into discrete codes. The underlying architecture, often rooted in the Generative Pre-trained Transformer (GPT) framework, such as GPT-3.5, forms the foundation for the model's ability to comprehend intricate musical elements. Through an intensive training process, the AI learns patterns associated with melody, harmony, and rhythm, culminating in a comprehensive understanding of musical composition. When tasked with music generation, users provide a seed input, such as a genre or mood, guiding the AI's creative output. The resultant compositions showcase the versatility of the model, offering original and coherent pieces across a spectrum of genres. Notably, some jukebox AI implementations allow for fine-tuning and customization, permitting users to influence specific elements such as tempo or instrumentation. In more recent developments, rudimentary singing was also implemented as well as AI "extensions" of a short user-inputted audio clip.

References

- Borsos, Zalán, Raphaël Marinier, Damien Vincent, Eugene Kharitonov, Olivier Pietquin, Matt Sharifi, Dominik Roblek, et al. 2023. “AudioLM: A Language Modeling Approach to Audio Generation.” *IEEE/ACM Transactions on Audio, Speech, and Language Processing* 31: 2523–33. <https://doi.org/10.1109/taslp.2023.3288409>.
- Engel, Jesse, Krishna Kumar, Shuo Agrawal, Ishaan Chen, Chris Gulrajani, Adam Donahue, and Roberts. n.d. “GANSYNTH: ADVERSARIAL NEURAL AUDIO SYNTHESIS.” <https://arxiv.org/pdf/1902.08710v2.pdf>.
- Van Den Oord, Aäron, Sander Dieleman, Heiga Zen, Karen Simonyan, Oriol Vinyals, Alex Graves, Nal Kalchbrenner, Andrew Senior, and Koray Kavukcuoglu. n.d. “WAVENET: A GENERATIVE MODEL for RAW AUDIO.” <https://arxiv.org/pdf/1609.03499v2.pdf>.